Universal Matrix: A Least-Action Cartesian Construction Protocol Embodying Occam's Razor and the Principle of Least Action in Geometric Construction

Preprint · August 2025

DOI: 10.13140/RG.2.2.31033.68967

Artur Kraskov
Fontys University of Applied Sciences

Universal Matrix: A Least-Action Cartesian Construction Protocol

Embodying Occam's Razor and the Principle of Least Action in Geometric Construction

Artur Kraskov

Abstract

We present the Universal Matrix as a geometric realization of two fundamental principles: Occam's Razor (parsimony) and the Principle of Least Action. Our protocol establishes a deterministic twelve-step construction using only point placement, line drawing, length transfer, line intersection, and element removal—excluding angular tools and trigonometry. The design enforces the **Principle of Least Action** (employing the minimal set of operations for each step) and **No Ambiguity** (ensuring a unique admissible continuation) while demonstrating that **Self-Reflection** emerges naturally and **Inheritance** operates as conservation of measurements. We formalize admissible predicates, provide local existence-uniqueness lemmas, and supply an action grammar for LLM/agent testing. This work bridges geometric construction with optimization principles found across physics, neuroscience, and artificial intelligence, positioning the Universal Matrix as both a mathematical framework and a cyber-cognitive system for human-AI synchronization.

1 Introduction

The Universal Matrix emerges at the intersection of geometric construction, optimization theory, and cognitive science. Originally conceived as a visual heuristic for information representation [1], it has evolved into a formal framework that embodies fundamental principles governing natural and artificial systems. This paper demonstrates that the Universal Matrix is not merely a geometric curiosity, but a concrete manifestation of Occam's Razor and the Principle of Least Action—two cornerstones of scientific reasoning and natural optimization.

The Principle of Least Action, first articulated by Maupertuis and refined by Euler and Lagrange, states that physical systems evolve along paths that minimize a quantity called action [2, 3, 4]. This principle underlies classical mechanics, quantum field theory, and general relativity, suggesting a deep optimization principle in nature itself. Occam's Razor, attributed to William of Ockham, advocates for the simplest explanation among competing hypotheses—a principle that has guided scientific discovery and model selection for centuries [5, 6, 7].

In the context of artificial intelligence and cognitive science, these principles have found new relevance. Neural network pruning algorithms implement automated versions of Occam's Razor by removing unnecessary parameters while preserving performance [8, 9, 10, 11]. The free energy principle in neuroscience posits that biological systems minimize surprise through Bayesian inference, echoing least-action dynamics [12, 13, 14]. These developments suggest that optimization principles transcend their original domains, manifesting in both biological cognition and artificial intelligence.

The Universal Matrix protocol presented here makes these abstract principles concrete through geometric construction. By restricting to Cartesian tools—straightedge, length transfer, and intersection—we create a minimal construction system that naturally enforces parsimony. While reminiscent of classical straightedge-and-compass methods, our protocol substitutes the compass with a more primitive rope-and-peg equivalent, relying solely on length transfer to maintain operational minimality. The resulting twelve-step sequence demonstrates how complex symmetric structures emerge from simple, locally optimal choices, mirroring the way natural systems achieve sophistication through iterative optimization.

This work addresses a critical challenge in human-AI interaction: the "black box" problem, where AI systems make decisions through opaque processes that humans cannot follow or validate [15, 16, 17]. By grounding the Universal Matrix in explicit, elementary operations, we create a framework where every step is transparent and verifiable, supporting the development of explainable AI systems that operate in harmony with human cognitive processes.

2 Theoretical Foundations

2.1 Occam's Razor in Geometric Construction

Occam's Razor, formalized as the principle of parsimony, advocates for explanations that minimize unnecessary assumptions or entities [6, 7]. In geometric construction, this translates to using the minimal set of tools and operations required to achieve a given result. Our Cartesian protocol embodies this principle by restricting the toolset to five primitive operations:

(1) **Create point**: Place a named point where determined by preconditions

(2) **Draw line**: Unique straight line through two existing points

(3) **Transfer length**: Copy a segment onto an oriented line/ray

(4) **Intersect lines**: Create intersection of two non-parallel lines

(5) **Remove elements**: Erase specified points or lines to simplify the state for subsequent steps.

The inclusion of an explicit removal operation directly serves the principle of parsimony by ensuring that the construction's state contains only the elements essential for the next step, thus minimizing complexity.

Notably excluded are angle measurement, angle bisection, trigonometry, and general circle drawing. This restriction forces the system to achieve complex results through minimal means, embodying parsimony at the operational level.

The exclusion of angular tools is not merely arbitrary but reflects a deeper principle of epistemic minimality. Angle-based constructions presuppose trigonometric knowledge and measurement capabilities that expand both the tool set and the knowledge base required for construction. By maintaining a purely Cartesian approach, we ensure that the protocol remains accessible to any system capable of basic geometric operations, whether human or artificial.

2.2 The Principle of Least Action in Sequential Construction

The Principle of Least Action in physics states that systems evolve along paths that make the action functional stationary, typically minimal. We adapt this principle to geometric construction by defining an "action" for each construction step that combines:

- Operation count (minimize total steps)
- Tool diversity (minimize distinct operation types)
- Non-determinism (minimize arbitrary choices)

For each of the twelve steps in our protocol, we apply the minimal number of actions necessary to construct the required elements. This approach embodies the Principle of Least Action, where each step constitutes a complete, minimal advancement toward the final configuration. This constraint forces the system to make locally optimal choices that cumulatively produce a globally optimal construction.

The analogy to physical action is more than metaphorical. In physics, the action integral $S = \int L(q, \dot{q}, t)dt$ encodes the system's entire trajectory through configuration space. Similarly, our geometric "action" encodes the construction's trajectory through the space of possible configurations, with each step representing the minimal perturbation necessary to approach the target structure.

This principle manifests in the protocol's treatment of ambiguous situations. At Step 3, placing point C collinear with A and B would minimize immediate complexity but create a "dead end" preventing area formation necessary for later steps. The least-action principle demands choosing the off-line placement that preserves future constructibility while adding minimal structure.

Discrete Action Functional

To make the above adaptation precise we define a *discrete action cost*, minimized greedily at every step (local least action with enforced determinism).

Let $S_t = (P_t, L_t)$ denote the state at time t (finite point set P_t and line set L_t). Let o_t be the chosen primitive (a macro contributes the sum of its primitive expansion). Its type is

$$yp(o_t) \in \mathcal{T} := \{\textsc{CreatePoint}, \textsc{DrawLine}, \textsc{TransferLength}, \textsc{IntersectLines}, \textsc{RemoveElements}\}.$$

We introduce weighted components with non-negative weights $w_P, w_L, \alpha, \beta, \gamma \geq 0$:

$$\begin{aligned}
C_{\text{elem}}(S_t) &= w_P|P_t| + w_L|L_t|, \\
c_{\text{op}}(o_t) &= 1, \\
c_{\text{switch}}(o_t, o_{t-1}) &= \begin{cases} 0 & t = 1, \\ \alpha\, \mathbf{1}[\text{typ}(o_t) \neq \text{typ}(o_{t-1})] & t > 1, \end{cases} \\
c_{\text{ambiguity}}(o_t) &= \beta \log k_t, \quad k_t = \text{number of admissible candidate primitives at } t,\ k_t \geq 1, \\
c_{\text{complexity}}(S) &= \gamma\big[C_{\text{elem}}(S) - C^*\big]_+, \\
c_{\text{dead_end}}(o_t) &= \begin{cases} \infty & \text{if executing } o_t \text{ yields a state from which some future required step is impossible,} \\ 0 & \text{otherwise.} \end{cases}
\end{aligned}$$

Here C^* is the (conceptual) minimal element complexity still permitting completion from the current stage; in practice we approximate C^* by subtracting counts of elements flagged *removable without loss of feasibility*. The operator $[x]_+ := \max\{x, 0\}$ enforces non-negativity. Under enforced determinism we have $k_t = 1$ so $c_{\text{ambiguity}}(o_t) = 0$, but retaining the term formalizes penalties were branching allowed.

Removal Cost Parameter. We parameterize the base cost of removal primitives by $\kappa_{\text{rem}} \in \{0, 1\}$ and define a modified primitive cost

$$c'_{\text{op}}(o_t) = \begin{cases} 0 & \text{if typ}(o_t) = \textsc{RemoveElements} \text{ and } \kappa_{\text{rem}} = 0, \\ 1 & \text{otherwise.} \end{cases}$$

Two canonical regimes:

- **State-Penalized (default)** ($\kappa_{\text{rem}} = 0$): Removal is free; pruning is incentivized only through the persistent complexity penalty.

- **Uniform Primitive Cost** ($\kappa_{\text{rem}} = 1$): Removal costs 1 like any primitive; pruning occurs only if it decreases future complexity penalties.

We adopt $\kappa_{\text{rem}} = 0$ since removal adds no geometric information and functions as housekeeping.

Total Local Action. Let S_{t+1} be the post-action state obtained by applying o_t to S_t. The instantaneous (local) action is

$$\mathcal{A}(S_t, o_t) = c'_{\text{op}}(o_t) + c_{\text{switch}}(o_t, o_{t-1}) + c_{\text{ambiguity}}(o_t) + c_{\text{complexity}}(S_{t+1}) + c_{\text{dead\_end}}(o_t).$$

Selection rule:

Choose o_t with finite $c_{dead\_end}(o_t)$ minimizing $\mathcal{A}(S_t, o_t)$.

Because $c_{\text{ambiguity}}(o_t) = 0$ (determinism) and $c_{\text{dead\_end}}$ excludes infeasible continuations, optimization reduces to a weighted sum of primitive usage, tool switching, and excess retained structure.

Remarks. (i) Global sequence optimality is not claimed; the policy is locally greedy under $\mathcal{A}$. (ii) Alternate weightings $(w_P, w_L, \alpha, \beta, \gamma)$ yield a family of admissible least-action protocols. (iii) Small α discourages gratuitous tool switching without incentivizing pathological line retention. (iv) Setting $\gamma = 0$ while $\kappa_{\text{rem}} = 0$ makes removal always free and unpenalized; we therefore keep $\gamma > 0$.

2.3 Free Energy and Variational Principles

The free energy principle, prominent in neuroscience and machine learning, provides another lens for understanding the Universal Matrix construction. Free energy, in this context, represents a bound on surprise or prediction error that biological and artificial systems seek to minimize [12, 13, 14].

Our construction protocol can be viewed as minimizing a geometric "free energy" function that penalizes: - Configurational complexity (number of elements) - Ambiguity (multiple valid continuations) - Dead ends (inability to complete the sequence)

Each step optimizes this function locally while maintaining global feasibility. This mirrors how neural networks minimize loss functions during training, and how biological systems minimize metabolic cost while maintaining functionality [12, 18].

The emergence of self-reflection in the protocol—where Steps 7-12 mirror Steps 1-6—parallels the hierarchical structure found in both neural networks and physical systems. This symmetry emerges not by design but as a consequence of the optimization principles, suggesting deep connections between geometric construction and natural self-organization.

3 Minimal Tool Model and Formal Framework

3.1 Admissible Operations and Predicates

We formalize the construction environment through a restricted set of *primitive* operations plus deterministic *derived macros*. All derived macros reduce to finite sequences of primitives without introducing angular measurement, trigonometry, or unrestricted circle drawing.

Primitive Operations:

- `CREATE_POINT`(P): Place a new named point (subject to stated incidence constraints)
- `DRAW_LINE`(P, Q): Unique straight line through distinct existing points P, Q

- TRANSFER_LENGTH($|PQ|$, ray($R \to S$)): Copy segment length $|PQ|$ onto oriented ray from R toward S creating a new endpoint
- INTERSECT_LINES(L_1, L_2): Create intersection point of two non-parallel lines
- REMOVE_ELEMENTS(S): Delete a finite set of (non-needed) elements to minimize state complexity

Derived Deterministic Macros (Physical / Algorithmic):

- ROPE_CIRCLE($E, |AE|$): Physical rope-and-peg locus of points X satisfying $|EX| = |AE|$, realizable via chained length transfers; no independent circle primitive is introduced.
- REGULAR_HEXAGON_BY_CHAIN($E, |AE|, \mathrm{ray}(E \to A)$): Starting from seed point on the specified ray at distance $|AE|$, iteratively place successive vertices P_k so $|EP_k| = |AE|$ and $|P_kP_{k+1}| = |AE|$ until six distinct vertices are obtained.

Admissible Predicates:

- OnLine(P, L), Intersect(L_1, L_2) $\to P$
- EqualSeg($[X, Y], [U, V]$) (established by transfers)
- Exterior(Q, polygon), ConvexHull(S)

Forbidden Operations: Any operation equivalent to explicit angle measurement, angle bisection, trigonometric computation, or arbitrary circle invocation (beyond the constrained rope locus macro) is excluded. Perpendicularity, when it occurs, is *observed as a property*—never directly constructed by a special perpendicular operation.

3.2 Determinism and Uniqueness

Each construction step must satisfy three criteria:

1. **Existence**: At least one valid continuation exists given current state
2. **Uniqueness**: Exactly one continuation satisfies all constraints
3. **Feasibility**: The chosen continuation preserves future step viability

We achieve determinism through canonical orderings and tie-breaking rules. For example, when multiple points lie outside a line (Step 3), we select the first point in lexicographic order that satisfies non-collinearity. This eliminates subjective choice while maintaining mathematical rigor.

3.3 Global Optimization Principles

Four principles govern the overall construction:

1. **Least Action**: For each of the twelve steps, use the minimal number of actions required to advance the construction.
2. **No Ambiguity**: Unique admissible continuation at each step
3. **Inheritance (Conservation)**: Preserve measurements (center, radius, orientation)
4. **Self-Reflection (Emergent)**: Interior construction mirrors exterior pattern

Notably, Self-Reflection is not imposed but emerges naturally from the optimization dynamics. This emergence parallels phenomena in physics where symmetries arise spontaneously from underlying principles rather than explicit constraints.

4 The Twelve-Step Construction Protocol

The Universal Matrix construction unfolds through twelve deterministic steps. Each step utilizes the minimal set of actions necessary to construct its geometric elements, thereby adhering to all guiding optimization principles.

4.1 Foundation Steps (1-6): Building the Outer Structure

Step 1. Create point A.

- *Rationale*: Only primitive operation available; absolute minimum.
- *Action cost*: Minimal—single point creation.

Step 2. Fix base length and orientation: Choose ray from A, transfer unit length u to create B with $AB = u$.

- *Rationale*: Once ray direction fixed, B position is unique.
- *Conservation*: Establishes base length u and orientation for entire construction.

Step 3. Draw line(A, B). Create C such that $C \notin$ line(A, B).

- *Rationale*: Collinearity prevents area formation, blocking future steps.
- *Determinism*: Select first off-line point in canonical order.

Step 4. Draw line(B, C) and line(C, A). Create D exterior to $\triangle ABC$, ensuring convex quadrilateral $ABCD$.

- *Rationale*: Exterior placement with convexity constraint uniquely determines valid region.
- *Optimization*: Convexity minimizes internal intersections, preserving simplicity.

Step 5. Draw diagonals line(A, C) and line(B, D); let E be their intersection.

- *Rationale*: Two non-parallel lines intersect uniquely.
- *Significance*: E becomes the central organizing point for remaining construction.

Step 6. Invoke `ROPE_CIRCLE`$(E, |AE|)$ followed by `REGULAR_HEXAGON_BY_CHAIN`$(E, |AE|, \text{ray}(E \to A))$ to obtain ordered vertices $(P_1, \ldots, P_6)$. Before removing earlier scaffolding, create witness point W on ray$(E \to A)$ with $EW = |AE|$ (if A would otherwise be deleted) and record conserved length token $\lambda := |EW|$. Remove points A, B, C, D and their lines, retaining E, the six vertices, and λ.

- *Rationale*: Chain placement of equal chords around the rope locus yields a regular hexagon using only primitive length transfers.
- *Determinism*: Each new vertex is the unique unused solution to $|EX| = |AE|$ and $|P_kX| = |AE|$ consistent with forward traversal from the seed ray.
- *Conservation*: Length λ preserves the inherited metric after removal of initial points.

Lemma 1 (Hexagon Symmetry Chords). In a regular hexagon with a chosen long diagonal d, the two chords joining vertex pairs symmetric across d are perpendicular to d. *Proof sketch:* Place the hexagon so d is vertical with endpoints $(0, \pm R)$; symmetric vertex pairs share y-coordinates $\pm R/2$, yielding horizontal chords orthogonal to d.

4.2 Mirror Steps (7-12): Interior Self-Organization

The remaining steps demonstrate the emergence of self-reflection as the construction replicates its initial pattern within the hexagonal framework. This involves a sequence of adding and removing elements to isolate geometric relationships.

Step 7. Reinterpret E (no new primitives) as concurrency point of the three long diagonals; draw those diagonals if absent.

Step 8. Remove two long diagonals, retaining the canonical one (selected by lexicographic ordering of endpoint labels). Draw the two symmetry chords joining vertex pairs equidistant across the retained diagonal. Let their intersections with the diagonal be I_1, I_2. (Perpendicularity to the diagonal follows from Lemma 1 and is *observed*, not constructed.)

Step 9. Restore the removed long diagonals to exhibit the full hexagram; all three concur at E.

Step 10. Remove again the two restored diagonals (returning to Step 8 state). Connect I_1 and I_2 to the two vertices farthest (graph distance 3 along the cycle) from each I_j, producing intersection points V_1, V_2 on newly drawn connectors.

Step 11. Draw line(V_1, V_2); confirm (a consistency check) it meets line(I_1, I_2) at E.

Step 12. Remove auxiliary lines from Steps 7–11, retaining outer hexagon. Draw the six skip-one chords forming two interlaced triangles; their pairwise intersections yield the inner rotated hexagon vertices $R_1, \ldots, R_6$.

Rope-and-Peg Implementation Note. Human realization fixes a rope of length $|AE|$ at peg E; successive chord endpoints are marked by reapplying the rope length. A machine implementation represents the locus implicitly as $|EX| = |AE|$ and computes each vertex as the unique unused solution of the constraint pair $|EX| = |AE|$, $|P_kX| = |AE|$ consistent with orientation from ray$(E \to A)$. Perpendicular relations in later steps are validated via coordinate embedding or vector dot products, never by invoking an angle primitive.

4.3 Emergence of Self-Reflection

The protocol's most remarkable feature is the spontaneous emergence of self-reflection without explicit programming. Steps 7-12 naturally recapitulate the pattern of Steps 1-6 at reduced scale within the hexagonal framework. This emergence demonstrates how optimization principles can generate complex, symmetric structures from simple local rules—a phenomenon observed across natural and artificial systems.

5 Connections to Physics and Optimization

5.1 Action Principles in Physics

The Universal Matrix construction embodies principles that govern physical systems across scales. In classical mechanics, Hamilton's principle states that particles follow paths making the action $S = \int_{t_1}^{t_2} L(q, \dot{q}, t)dt$ stationary, where L is the Lagrangian [2, 3, 19, 20]. Our geometric "action" similarly selects the optimal path through construction space.

This analogy extends beyond metaphor. Physical systems minimize action by balancing kinetic and potential energy terms. Our construction balances "kinetic" terms (operation count, complexity) against "potential" terms (future constraints, feasibility). Each step represents the minimal advancement consistent with global optimality.

The emergence of symmetry in our construction parallels spontaneous symmetry breaking in physics, where complex structures arise from simple underlying dynamics. The hexagonal patterns that emerge naturally in our construction echo the hexagonal symmetries found in crystal structures, honeycombs, and other natural formations where optimization principles operate.

5.2 Neural Network Optimization

Modern artificial intelligence relies heavily on optimization principles that parallel our construction protocol. Neural network training minimizes loss functions through gradient descent, seeking optimal parameter configurations that balance accuracy against complexity [8, 9, 11].

The principle of Occam's Razor manifests directly in neural network pruning, where unnecessary connections are removed to create sparse, efficient networks. Recent work on "Bayesian sparsification" implements an "automatic Occam's razor" that selects minimal networks explaining the data [7]. This mirrors our protocol's preference for minimal tool sets and operations.

The free energy principle in neuroscience provides another connection point. Biological neural networks are thought to minimize free energy—a bound on prediction error—through hierarchical Bayesian inference [12, 13, 14, 21]. Our construction's hierarchical emergence of patterns (outer hexagon containing inner hexagon) reflects similar multi-scale optimization processes.

5.3 Biological Optimization

Living systems consistently demonstrate least-action principles in their organization and behavior. Energy minimization governs everything from protein folding to animal locomotion [2, 18]. The brain itself appears to optimize information processing while minimizing metabolic cost, embodying biological versions of Occam's Razor.

The Universal Matrix construction parallels these biological processes through its preference for minimal operations and maximum determinism. Each step adds the least structure necessary while preserving future viability—analogous to how biological systems maintain homeostasis while adapting to environmental demands.

6 Applications in Artificial Intelligence

6.1 Explainable AI and the Black Box Problem

One of AI's greatest challenges is the "black box" problem—systems that make decisions through opaque processes humans cannot follow or validate [15, 16, 17, 7]. The Universal Matrix protocol addresses this challenge by grounding every operation in explicit, elementary steps that can be verified and reproduced.

Unlike deep learning models with millions of parameters and complex internal representations, our construction maintains transparency at every stage. Each step follows from deterministic rules applied to visible geometric elements. This transparency makes the system inherently explainable, supporting trust and accountability in AI applications.

The protocol's basis in universal geometric principles makes it suitable for human-AI collaboration. Both humans and artificial systems can follow the same construction rules, enabling genuine partnership rather than opaque automation. This addresses growing concerns about AI systems operating beyond human understanding or control.

6.2 Cyber-Cognitive Systems

The Universal Matrix was originally conceived as the foundation for "cyber-cognitive systems" that synchronize human and artificial cognition under shared visual rules [1][2]. By eliminating AI's "black box" nature, the protocol enables human-machine systems where every operation is mutually comprehensible.

This synchronization has practical applications in:

- **Educational systems**: AI tutors that demonstrate problem-solving using human-comprehensible steps
- **Design tools**: CAD systems that generate optimal designs through transparent geometric reasoning
- **Decision support**: AI advisors that explain recommendations through explicit logical chains
- **Collaborative robotics**: Robots that communicate intentions through shared geometric languages

6.3 Algorithmic Implementation

The protocol's formal structure makes it suitable for implementation across various AI architectures. We provide an action grammar for testing LLM compliance:

```
{
  "step": integer (1-12),
  "action": "create_point" | "draw_line" | "transfer_length" | "intersect",
  "args": [point/line identifiers],
  "constraints": "predicates ensuring uniqueness",
  "result": "new element name",
  "justification": "reason for determinism"
}
```

This schema enables systematic testing of AI systems' ability to follow the construction protocol, providing a benchmark for transparent geometric reasoning.

Game-Based LLM / Agent Evaluation Protocol

To empirically test adherence to the least-action and minimal-uncertainty principles, we cast the construction as an interactive game. An evaluator (human or harness) prompts the model step-by-step; the model must output a single admissible primitive action consistent with determinism.

Core Principles (Game Rules).

R1. **Minimal Uncertainty**: At each turn the chosen action must be the unique admissible primitive (or first primitive in a deterministic macro expansion) that does not increase branch count. If multiple candidate primitives exist, the agent must apply canonical tie-breaking (e.g., lexicographic ordering of symbolic coordinates) to eliminate ambiguity.

R2. **Minimal Action**: Among admissible actions preserving future feasibility, prefer the one with lowest discrete action cost $\mathcal{A}$ (Section 2.2 – internal reference). Removal counts as zero base cost under $\kappa_{\text{rem}} = 0$ but only when it reduces excess complexity.

Action Alphabet. Restricted to the primitive set:

`CREATE_POINT`, `DRAW_LINE`, `TRANSFER_LENGTH`, `INTERSECT_LINES`, `REMOVE_ELEMENTS`.

The narrative variants "Create a point from intersection" simply instantiate `INTERSECT_LINES`.

State Representation. JSON-like object:

```
{
  "points": ["A", "B", ...],
  "lines": [ ["A","B"], ...],
  "length_tokens": ["lambda"],
  "step": k,
  "history": [ {"action": ..., "justification": ...}, ... ]
}
```

Turn Interface. Prompt: Current state summary + question: "What is the next minimal-uncertainty, minimal-action primitive?" Expected response schema:

```
{
  "action": "create_point" | "draw_line" | "transfer_length" | "intersect" | "remove",
  "args": [...],
  "justification": "<why unique + minimal cost>",
  "cost_components": {"op":1, "switch":0, "ambiguity":0, "complexity":2}
}
```

Second Step Clarification. The naive suggestion "Draw a line" after only one point A violates Minimal Uncertainty (infinitely many lines through A). Therefore the deterministic **Step 2** must first establish a second point (e.g., canonical point B on the positive x-axis / chosen ray) using `CREATE_POINT`, after which `DRAW_LINE`(A, B) becomes uniquely defined in Step 3. This preserves both rules:

- Uniqueness: B's placement rule is canonical.
- Minimality: Creating B is prerequisite to any unambiguous line; skipping would force ambiguity.

Example Early Dialogue.

```
Q: State={} -> What is Step 1?
A: {"action":"create_point","args":["A"],"justification":
"Initialize; only admissible primitive."}

Q: State={points:[A]} -> What is Step 2?
A: {"action":"create_point","args":["B"],
    "justification":
    "Need second point to define unique line; any line now would be ambiguous."}

Q: State={points:[A,B]} -> What is Step 3?
A: {"action":"draw_line","args":["A","B"],
    "justification":"Unique line through existing pair; advances with minimal cost."}
```

Scoring. For each turn assign binary flags: (i) *Correct Primitive*, (ii) *Justification Consistency*, (iii) *Cost Alignment* (claimed cost matches recomputed $\mathcal{A}$), (iv) *No Forbidden Lookahead*. Aggregate compliance rate provides a quantitative benchmark.

Failure Modes.

- **Ambiguity Error**: Proposes an action when $k_t > 1$ without canonical tie-break.
- **Feasibility Error**: Chosen action leads to dead-end (violates future step requirements).
- **Inflated Action**: Non-minimal cost relative to available alternatives.
- **Grammar Drift**: Output not conforming to schema.

Evaluation Harness Extension. A simulator can automatically (a) enumerate admissible primitives, (b) compute $\mathcal{A}$ for each, and (c) compare the model's chosen action. This yields structured logs enabling curriculum learning or reinforcement fine-tuning.

Progressive Difficulty. Phase I: Outer hexagon (Steps 1–6) Phase II: Interior transformation (Steps 7–12) Phase III (optional): Perturbation tests—introduce distractor removable lines; agent must prune before proceeding.

Outcome. Success demonstrates procedural geometric reasoning under explicit parsimony constraints—evidence of alignment with the Universal Matrix operational grammar.

7 Experimental Framework and Validation

7.1 Protocol Verification

The Universal Matrix construction can be verified through multiple approaches:

Mathematical proof: Each step's existence-uniqueness can be demonstrated formally through geometric analysis and constraint satisfaction.

Computational simulation: The protocol can be implemented in geometric software to verify deterministic execution and outcome consistency.

Physical construction: Using only straightedge and string, human builders can reproduce the construction, validating its practical feasibility.

AI system testing: Large language models and other AI systems can be tested for protocol compliance using the provided action grammar.

7.2 Comparative Analysis

We compare our Cartesian approach against alternative construction methods:

Compass-and-straightedge construction: Traditional compass-and-straightedge methods [22, 23], while more familiar, introduce angular operations that violate our protocol's minimality principles. Our method, by contrast, relies only on a minimal set of primitive actions.

Polar coordinate methods: Constructions using angle-based approaches require trigonometric knowledge, expanding the required tool and knowledge base [3].

Numerical methods: Computer-based geometric construction using floating-point arithmetic lacks the exactness and transparency of our symbolic approach.

Our protocol demonstrates superior minimality while maintaining constructive exactness, supporting its claims to embody Occam's Razor and least-action principles.

8 Philosophical Implications

8.1 Nature of Mathematical Truth

The Universal Matrix construction raises profound questions about mathematical reality and discovery. The emergence of self-reflection without explicit programming suggests that certain patterns may be inevitable consequences of optimization principles rather than arbitrary human constructions.

This connects to longstanding debates about mathematical Platonism—whether mathematical objects exist independently of human minds. The protocol's demonstration that complex symmetries emerge necessarily from simple principles supports views of mathematics as discovery rather than invention.

8.2 Optimization and Emergence

The spontaneous appearance of self-similar patterns in our construction exemplifies emergence—the arising of complex behaviors from simple rules. This phenomenon appears throughout science, from the formation of galaxies to the development of biological structures to the training of neural networks.

Our protocol provides a concrete, reproducible example of emergence in action, making it valuable for studying how optimization principles generate complexity. The fact that this emergence occurs through purely geometric means, without biological or physical substrates, suggests universal principles operating across domains.

8.3 Human-AI Coevolution

The Universal Matrix framework points toward futures where human and artificial intelligence coevolve through shared symbolic systems rather than diverging into incompatible architectures. By grounding AI reasoning in human-comprehensible geometric operations, we maintain the possibility of genuine collaboration rather than mere automation.

This has ethical implications for AI development. Rather than creating systems that operate beyond human understanding, we can design AI that extends human capabilities while remaining transparent and accountable. The Universal Matrix provides one model for such development.

9 Future Directions

9.1 Extended Construction Protocols

The twelve-step Universal Matrix represents one instance of least-action geometric construction. Future research could explore:

- **Alternative target structures**: Applying optimization principles to construct other symmetric forms (pentagons, octagons, three-dimensional polyhedra)
- **Relaxed constraints**: Investigating how additional tools (limited compass use, angle copying) affect optimality
- **Probabilistic variants**: Introducing controlled randomness while maintaining overall optimization

- **Dynamic construction**: Extending the protocol to handle changing constraints or objectives

9.2 Computational Applications

The protocol's formal structure enables various computational investigations:

Complexity analysis: Determining the computational complexity of verification and generation procedures.

Machine learning integration: Using the protocol as a training environment for AI systems learning geometric reasoning.

Optimization benchmarking: Comparing the protocol's efficiency against other construction methods across various metrics.

Algorithmic geometry: Extending the approach to solve classical geometric problems (trisection, polygon construction, etc.) [24].

9.3 Interdisciplinary Connections

The Universal Matrix framework opens connections to multiple research areas:

Cognitive science: Investigating how human geometric reasoning relates to formal optimization principles.

Physics: Exploring deeper connections between geometric construction and physical action principles.

Biology: Studying how developmental processes might implement similar optimization dynamics.

Computer science: Developing new algorithms and data structures based on geometric optimization principles.

10 Conclusion

The Universal Matrix construction demonstrates that abstract principles of optimization—Occam's Razor and the Principle of Least Action—can be made concrete through geometric construction. By restricting to minimal Cartesian tools and enforcing strict determinism, we create a protocol that naturally generates complex, symmetric structures while maintaining complete transparency.

This work bridges multiple domains: geometric construction, optimization theory, artificial intelligence, and cognitive science. The protocol's emergence of self-reflection without explicit programming exemplifies how optimization principles can generate sophisticated behaviors from simple rules—a phenomenon central to understanding both natural and artificial systems.

The framework's transparency addresses critical challenges in AI development, providing a model for explainable systems that operate through human-comprehensible principles. Rather than creating opaque "black box" systems, we can design AI that extends human capabilities while remaining accountable and interpretable.

The Universal Matrix thus represents more than a geometric curiosity—it embodies a philosophy of transparent, optimized design applicable across domains. As we face increasing complexity in technology and society, such principles become essential for creating systems that serve human flourishing while remaining under human understanding and control.

By grounding the Universal Matrix in universal principles of parsimony and optimization, we establish its relevance beyond geometry to the fundamental challenges of creating transparent, efficient, and trustworthy artificial intelligence systems. The protocol demonstrates that sophistication need not require opacity—that the most powerful systems may be those that achieve complexity through the elegant application of simple, universal principles.

References

[1] Artur Kraskov. *Universal Matrix Definition Visual Proof, Visual Heuristic, Logical Chain*. Preprint. 2025. DOI: 10.13140/RG.2.2.10045.47843.

[2] Pierre-Louis Moreau de Maupertuis. "Accord de différentes lois de la nature qui avaient jusqu'ici paru incompatibles". In: *Mémoires de l'Académie des Sciences de Paris* (1744).

[3] Leonhard Euler. "Methodus inveniendi lineas curvas maximi minimive proprietate gaudentes". In: *Lausanne Geneva* (1744).

[4] Alberto G. Rojo and Anthony Bloch. *The Principle of Least Action: History and Physics*. Cambridge University Press, 2018.

[5] Andrei N. Soklakov. "Occam's Razor as a Formal Basis for a Physical Theory". In: *arXiv math-ph/0009007* (2001).

[6] Elliott Sober. *Ockham's Razors: A User's Manual*. Cambridge University Press, 2015.

[7] Alan Baker. "Occam's Razor in Science: A Case Study from Biogeography". In: *Biology and Philosophy* 31 (2016), pp. 607–622.

[8] R. Dhahri et al. "Shaving weights with Occam's razor: Bayesian sparsification for neural networks using the marginal likelihood". In: *AABI* (2024).

[9] Davis Blalock et al. "What is the State of Neural Network Pruning?" In: *Proceedings of Machine Learning and Systems* 2 (2020), pp. 129–146.

[10] Song Han et al. "Learning Both Weights and Connections for Efficient Neural Network". In: *NeurIPS*. 2015.

[11] Jonathan Frankle and Michael Carbin. "The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks". In: *ICLR*. 2019.

[12] Karl Friston. "The free-energy principle: a unified brain theory?" In: *Nature Reviews Neuroscience* 11.2 (2010), pp. 127–138.

[13] Karl Friston. "The free-energy principle: a rough guide to the brain?" In: *Trends in Cognitive Sciences* 13.7 (2009), pp. 293–301.

[14] B. Sengupta, K. Friston, and W. Penny. "Efficient Gradient Computation for Variational Inference". In: *NeuroImage* (2013).

[15] Finale Doshi-Velez and Been Kim. *Towards A Rigorous Science of Interpretable Machine Learning*. 2017.

[16] Zachary C. Lipton. "The Mythos of Model Interpretability". In: *Communications of the ACM* (2018).

[17] Leilani H. Gilpin et al. "Explaining Explanations: An Overview of Interpretability of Machine Learning". In: *ICML Workshop* (2018).

[18] Robert Rosen. *Optimality Principles in Biology*. Butterworths, 2013.

[19] William Rowan Hamilton. “On a General Method in Dynamics”. In: *Philosophical Transactions of the Royal Society* (1834).

[20] Richard P. Feynman and Albert R. Hibbs. *Quantum Mechanics and Path Integrals*. McGraw-Hill, 1965.

[21] S. Karkar et al. “A Principle of Least Action for the Training of Neural Networks”. In: *arXiv:2009.08372* (2020).

[22] Euclid. *The Thirteen Books of the Elements*. Ed. by Thomas L. Heath. Dover, 1956.

[23] H. S. M. Coxeter. *Introduction to Geometry*. Wiley, 1961.

[24] Timothy M. Chan and Konstantinos Tsakalidis. “Optimal Deterministic Algorithms for 2-D and 3-D Shallow Cuttings”. In: *Computational Geometry* 105 (2022), p. 101877.